

PostgreSQL-Backed Transport Analytics for Paris and New York City

Maksym DOLHOV

Mehdi AGHAEI

Ho Bao Khanh Nguyen

Nima DAVARI

April 13, 2026

Abstract

This project studies public-transport demand in Paris and New York City through a PostgreSQL-backed analytical workflow designed to keep the data source, derived views, Python analysis, notebook, and final paper aligned. Raw source tables remain in the `public` schema, while reporting-clean views in the `transport` schema define the contract consumed by the Python workflow. The final method set contains four components: temporal profiling, lag-based forecasting, anomaly detection, and contributor / city-structure comparison. Using the current cleaned reporting scope, the workflow produces 12,696 daily observations and 2,003 contributor-level rows. The forecasting baseline reaches an overall mean absolute percentage error of 6.02%, with lower error for NYC (5.86%) than Paris (6.94%). The anomaly layer flags a substantially higher rate in Paris (8.96%) than in NYC (0.73%). Paris also shows higher average daily demand (≈ 4.13 million) than NYC (≈ 2.63 million), while NYC displays a weaker weekday concentration. The resulting paper emphasizes reproducibility, reporting consistency, and defensible interpretation rather than a large model stack.

1 Introduction

Urban transit systems generate dense operational traces that can support demand monitoring, comparative planning, and baseline forecasting. Smart-card and ridership logs are especially useful because they preserve temporal structure at station or network level while remaining close to the decisions that transport operators need to make. For a student project, however, there is a second problem beyond the analysis itself: the workflow has to stay reproducible when multiple people work on notebooks, SQL, figures, and report text at the same time.

This project compares Paris and New York City transport demand through one shared analytical contract. Earlier repository iterations mixed local files, exploratory notebook material, and placeholder report claims. The final baseline removes that ambiguity. PostgreSQL stores the authoritative data, SQL views expose reporting-clean inputs, Python generates result tables and figures, and the report is written from those generated artifacts instead of from temporary notebook output. This matters because the scientific value of the project depends not only on the metrics that were produced, but also on the ability to trace each claim back to one stable result file.

The contribution of this final stage is therefore twofold. First, it produces a comparative demand analysis for Paris

and NYC using four methods: temporal profiling, lag-based forecasting, anomaly detection, and contributor / city-structure analysis. Second, it frames these methods inside a reproducible PostgreSQL-first workflow that keeps the final paper, notebook, and figures numerically aligned. The report focuses on practical questions: how the two systems differ over time, how well a simple forecasting baseline performs, which observations are anomalous, and which stations dominate total validated or observed demand.

The report is intentionally conservative in scope. It does not claim cross-city causal equivalence or state-of-the-art predictive performance. Instead, it shows that a medium-scale transport analysis can still be scientifically useful when the workflow is explicit, the result files are stable, and every major claim is traceable to a generated artifact. That is why the architecture discussion appears as part of the methodology rather than as separate internal engineering documentation.

2 Related Work

Smart-card data has long been used to recover travel behavior, rider segments, and network-level demand structure. Ma et al. [1] showed that transaction records can reveal stable travel patterns without relying on traditional survey frequency, while Long and Thill [2] combined smart-card data with household survey information to link ridership behavior with broader urban structure. Cats [3] further illustrated how smart-card traces can be mined for recurring mobility patterns at scale. These studies motivate the descriptive side of the current project, especially the use of station-level and city-level summaries instead of a single aggregate demand series.

Recent work also emphasizes richer behavior discovery. Aminpour and Saidi [4] used topic-modeling ideas to move beyond simple home-work simplifications, while Sun et al. [5] described how network analysis and smart-card records can expose structurally important stations. That literature is directly relevant to the contributor ranking used here: even when modeling remains simple, the rank order of highly used stations provides meaningful evidence about concentration, interchange hubs, and asymmetry between systems.

Forecasting literature is substantially more ambitious than the baseline used in this project. Dynamic graph and spatio-temporal learning approaches such as the model of Xie et al. [6] aim to capture inter-station dependencies directly, while ASTIR [7] demonstrates how spatio-temporal data mining can support crowd-flow prediction. Those

models are more expressive than the lag-based random forest baseline used here, but they also require more infrastructure and stronger assumptions about graph structure, feature availability, and tuning. For the final paper, the baseline choice is deliberate: the objective is to establish a transparent and reproducible comparison before moving to heavier architectures.

Anomaly and disruption monitoring provide another important thread. Kolios et al. [8] studied monitoring recurrent mobility patterns through adaptive event-triggering logic, which is conceptually close to the goal of identifying observations that diverge strongly from routine behavior. The present project does not attempt causal attribution, but it uses anomaly scoring as an analytical screening layer that helps highlight periods worth interpretation in the report.

Other recent papers are useful mainly as boundary markers for what this project does *not* attempt. He et al. [9] analyzed factors that influence station ridership in a specific metro context, while Chen et al. [10] explored trip-purpose and socio-economic inference from transit users. Kundu et al. [11] examined heterogeneous effects from a new transit line. Those studies show how quickly transport analytics can become policy-rich and context-heavy. The present paper remains more modest: it uses the literature to justify its analytical questions, but it stops short of causal or socio-economic interpretation that the current dataset cannot support directly.

Taken together, the literature supports the structure of this paper. Existing research shows that smart-card and ridership traces are suitable for descriptive mining, forecasting, and disruption-aware analysis. The distinct contribution here is not a new predictive architecture; it is the combination of a reporting-clean PostgreSQL contract, a reproducible Paris-vs-NYC comparison, and a final result set whose metrics, figures, and narrative are all generated from the same workflow outputs. The project therefore occupies a practical position between exploratory course work and a fully specialized transport-modeling study.

3 Data Description

The final project uses three operational transport datasets already loaded into PostgreSQL, together with two metadata tables. Paris demand is represented by `public.idfm_daily_validations`, which contains station-level daily validation counts, and `public.idfm_hourly_profiles`, which contains hourly distribution shares used for temporal shape analysis. NYC demand is represented by `public.mta_hourly_ridership`, an hourly ridership table that is later aggregated into comparable daily demand views. Reference tables `public.idfm` and `public.mta` provide descriptive network metadata such as station and network names.

The two cities are not perfectly symmetric in source structure. Paris provides direct daily station validations plus separate hourly profile shares, while NYC starts from hourly ridership observations that must be rolled up into daily station totals. The project therefore does not claim

Table 1: Primary data sources used in the final paper

Source table	City	Temporal grain	Reporting coverage
IDFM daily validations	Paris	daily station validations	2015–2024
IDFM hourly profiles	Paris	hourly profile shares	cleaned reporting years
MTA hourly ridership	NYC	hourly station ridership	2020–2024

one-to-one measurement equivalence at every intermediate step. Instead, it defines a cleaned reporting scope where both systems can be compared at daily city and station level while preserving the source-specific hourly views needed for temporal interpretation.

The cleaned reporting scope excludes incomplete boundary years from the final comparison-oriented views. In practice, that removes the partial Paris 2025 year from the analytical contract used by the report. The resulting output set contains 12,696 daily rows and 2,003 contributor-level station rows according to `analysis_summary.json`. These totals are large enough to support descriptive and baseline predictive analysis while still being small enough to keep the project workflow understandable and reproducible.

Two additional data choices are important. First, the final contributor analysis uses station-grain Paris rows only, which avoids mixing line-level and station-level semantics in the ranking output. Second, the cleaned reporting views apply the same year-scope logic across daily, contributor, and city-structure outputs. This consistency matters because it prevents the paper from comparing annual summaries built from one scope with contributor totals built from another.

The generated result files used by the paper reflect that same scope. Temporal outputs are written to separate yearly, monthly, weekday, and hourly CSV files; forecasting outputs are split into overall metrics, city metrics, and holdout predictions; anomaly outputs are split into row-level flags and city-level rates; and contributor outputs are separated from city structure summaries. This file-level separation makes it possible to trace each figure and table in the paper back to one concrete artifact rather than to a mixed notebook state.

Comparability therefore has to be handled procedurally rather than assumed. The paper does not argue that a Paris validation count is physically identical to an NYC ridership count at every level of aggregation. What it argues is narrower: once both systems are mapped into cleaned daily demand views with stable city labels, reporting scope, and contributor grain, they can support a defensible comparative analysis of trend, forecastability, anomaly concentration, and hub dominance.

That distinction is important for scientific writing. A project can be honest about source heterogeneity and still

PostgreSQL-First Transport Analytics Flow



Figure 1: PostgreSQL-first workflow used for the final generated outputs.

produce useful comparative evidence if the alignment rules are explicit. The PostgreSQL contract serves exactly that role here. It documents what is aligned, what remains source-specific, and which outputs should be treated as directly comparable in the final report.

4 PostgreSQL-Centered Architecture

PostgreSQL was chosen as the analytical backbone because the project needed one stable source of truth that multiple artifacts could consume without reloading or cleaning raw files separately. A CSV-first workflow would have left the report, notebook, and scripts vulnerable to local path drift and inconsistent intermediate cleaning steps. By contrast, PostgreSQL allows the raw uploaded tables to remain intact in the `public` schema while the cleaned reporting contract is expressed through SQL views in a separate `transport` schema.

That schema split is central to the final design. The `public` schema holds source-level operational tables. The `transport` schema holds reporting-ready views such as `v_daily_demand`, `v_nyc_hourly_profile`, `v_paris_hourly_profile`, `v_contributor_source`, and `v_city_structure_source`. This means the SQL layer, not the notebook, defines the cleaned year scope and contributor label normalization used in the final report. The Python workflow then consumes those views and writes the derived CSV outputs under `report/results/`, which are in turn used by the notebook, figure builder, and paper text.

This design improves reproducibility in two ways. First, it removes dependence on local raw files from the official workflow path. Second, it makes each quantitative claim traceable to either a SQL view or a generated result file. For a project of this size, that tradeoff is more valuable than adopting a more complex deployment architecture. The important design decision is therefore the analytical contract, not the underlying server hardware, which is intentionally outside the scope of this paper.

In practice, the contract works as follows. SQL views first standardize date scope, contributor labels, and daily aggregation. The Python workflow then constructs time-series features and writes the final CSV outputs. The notebook and report consume those outputs rather than issuing independent analytical queries. That separation reduces

Table 2: Key PostgreSQL analytical views

View	Role in the final workflow
<code>v_daily_demand</code>	Cleaned daily demand baseline for downstream methods.
<code>v_nyc_hourly_profile</code>	NYC hourly ridership shape for temporal interpretation.
<code>v_paris_hourly_profile</code>	Paris hourly validation-share shape for temporal interpretation.
<code>v_contributor_source</code>	Station-grain contributor totals for ranking.
<code>v_city_structure_source</code>	City-level aggregates for structural comparison.

the chance that the same project tells slightly different numerical stories in different artifacts, which was a real risk in the earlier local-file version of the repository.

5 Methods

5.1 Temporal profiling

Temporal profiling is the descriptive baseline for the whole project. The daily demand view is summarized by year, month, and weekday to reveal long-horizon trends, seasonal structure, and weekly rhythm. The two hourly profile views are kept separate because Paris hourly data is based on validation shares, while NYC hourly data is based on observed ridership totals. This method produces the tables `temporal_yearly_totals.csv`, `temporal_monthly_profile.csv`, `temporal_weekday_profile.csv`, `temporal_nyc_hourly_profile.csv`, and `temporal_paris_hourly_profile.csv`. The goal is not only to describe the two systems, but also to establish whether later forecasting and anomaly results are plausible relative to the seasonal structure visible in the data. Because this method operates directly on the cleaned reporting layer, it also serves as a validation step for the SQL contract: if the yearly totals or monthly patterns looked implausible, the rest of the workflow would immediately be suspect.

5.2 Lag-based forecasting

The forecasting stage uses the cleaned daily demand series as input. The model is intentionally simple: a random-forest regressor trained on lag features, rolling-window statistics, and calendar attributes derived from the daily series. The numerical features are `lag_1`, `lag_7`, `lag_14`, rolling means and standard deviations over 7-day and 30-day windows, and one-step percentage change. Calendar features include day of week, week of year, day of year, month, year, and weekend status. Region, source, and metric type are one-hot encoded as categorical inputs. The split is time ordered, not random, so the model is evaluated on future observations relative to the training period. The random forest uses 220 trees with `min_samples_leaf=2`, which balances stability and transparency for the project scale. This produces `forecast_metrics_overall.csv`,

Table 3: Method-to-output mapping in the final workflow

Method	Core inputs	Generated outputs
Temporal profiling	cleaned daily demand, hourly profile views	yearly, monthly, weekday, and hourly summaries
Forecasting	lag features, rolling features, calendar features	forecast metrics and holdout predictions
Anomaly detection	daily demand, lag features, rolling z-scores	anomaly flags and anomaly-rate summary
Contributors / structure	station-level facts, city daily aggregates	top contributors and structural comparison tables

forecast_metrics_by_city.csv, and forecast_predictions.csv. The objective is to provide a transparent baseline whose errors are easy to interpret rather than to optimize predictive performance at all costs.

5.3 Anomaly detection

The anomaly layer combines two signals: a rolling z-score that captures local deviation from recent behavior, and an IsolationForest model fitted on current value and lag-derived features. The anomaly feature set includes current value, 1-day and 7-day lags, a 7-day rolling mean and standard deviation, one-step percentage change, and a 30-day rolling z-score. The isolation model uses a contamination value of 0.03, while the statistical branch flags any observation with $|z| \geq 2.5$. This hybrid design helps separate routine daily variation from observations that are unusual both in local context and in the broader feature space. The output is written to anomaly_flags.csv and summarized in anomaly_rates.csv. The method is descriptive: it highlights unusual observations for interpretation, but it does not claim to identify the causal mechanism behind each anomaly.

5.4 Contributor and city-structure analysis

The final method ranks stations by total demand contribution and summarizes each city through aggregate statistics such as average daily demand, variability, weekend sensitivity, and peak month. Paris contributor analysis is explicitly restricted to station-grain rows to avoid mixing station and line semantics in the final ranking. NYC contributor totals are computed from the cleaned station daily demand layer. This stage produces top_contributors.csv, city_structure_summary.csv, and city_monthly_structure.csv. It complements the other methods by showing whether city-level differences are driven by broad network structure, dominant hubs, or both.




Figure 2: Current temporal outputs: monthly demand patterns and the NYC hourly profile.

5.5 Evaluation protocol

The evaluation logic follows the goal of the project rather than a benchmark competition. Forecasting is judged through holdout error on a time-ordered split because random shuffling would leak future information into the training set. City-level metrics are reported alongside overall metrics to avoid hiding cross-city asymmetry inside one aggregate score. For anomaly detection, the key reported output is not a classification accuracy measure, since there is no ground-truth anomaly label set in the project data. Instead, the paper reports anomaly rates and interprets the flagged periods qualitatively. For contributor and city-structure analysis, plausibility of station rankings and consistency with temporal summaries are treated as important validity checks.

This evaluation design fits the scope of the project. The goal is not to claim optimality for one algorithm, but to produce a final report whose methods are transparent enough to defend and whose outputs can be regenerated. In that sense, evaluation is tied as much to interpretability and consistency as to predictive performance alone.

The four-method combination was chosen for the same reason. Temporal profiling anchors the descriptive structure of the data. Forecasting tests whether that structure is strong enough to support short-term prediction. Anomaly detection checks where routine assumptions break down. Contributor and city-structure analysis then reconnect the time-series view to physically meaningful stations and system-level concentration. Each method answers a different question, but together they create a more complete analytical narrative than any one of them would alone.

6 Results

6.1 Temporal patterns

The yearly totals confirm the effect of the reporting-clean scope. NYC covers 2020–2024, while Paris covers 2015–2024 with the incomplete 2025 year excluded. The Paris series shows a larger absolute scale throughout most of the window and a visible collapse in 2020 before a multi-year recovery. NYC shows a similar pandemic-era break, but the amplitude is smaller in absolute terms because the baseline level is lower.

Monthly structure also differs. The generated monthly profile indicates that NYC peaks in February in the cleaned result set, while Paris peaks in October. Paris weakens sharply in August, which is consistent with seasonal travel

Table 4: Selected yearly totals from the cleaned reporting scope

City	2020	2023	2024
NYC	643.3M	1,163.3M	1,211.5M
Paris	926.0M	1,729.7M	1,024.2M

Table 5: Forecast accuracy by city

City	Rows	MAE	RMSE	MAPE
NYC	3640	29,402	62,965	5.86%
Paris	636	187,446	376,361	6.94%

behavior and reduced urban activity during summer holidays. At the weekly level, both systems show reduced weekend demand, but the later city-structure results confirm that the Paris weekend drop is stronger in relative terms.

The temporal outputs also provide a useful sanity check for the forecasting task. The series are not stationary in a strict sense: both cities contain a pandemic-era structural break, strong seasonality, and pronounced weekly effects. A forecasting baseline that achieves low error in spite of that non-stationarity is therefore capturing real structure rather than merely memorizing a flat process.

The selected yearly totals in Table 4 make the broad recovery story visible without requiring the reader to parse the full CSV output. Both systems rebound strongly after 2020, but the size and persistence of the recovery differ. NYC climbs steadily through 2024 in the cleaned result set, whereas Paris rebounds sharply by 2023 and then drops in the 2024 total present in the current reporting output. That pattern is part of the generated baseline and should be discussed as an observed result rather than normalized away by the narrative.

6.2 Forecasting

The forecasting baseline is strong enough to support a substantive results section without overstating the model. According to `forecast_metrics_overall.csv`, the final split trains on 8,336 rows and evaluates on 4,276 rows. The overall error reaches an MAE of 52,909, RMSE of 156,343, and MAPE of 6.02%. These values show that a simple lag-based baseline can already explain a large part of the daily structure in the cleaned reporting series.

The city-level split in Table 5 reveals an important asymmetry. NYC produces a lower MAE, lower RMSE, and lower MAPE than Paris. The most conservative reading is that the NYC series is easier to predict with the current feature set, while Paris contains stronger short-term swings or structure that the baseline model does not fully capture. The gap is not catastrophic: both cities still remain within a single-digit MAPE range. That is a useful project finding because it supports the report narrative without requiring a more opaque forecasting stack.

The absolute Paris error is much larger because the Paris city series operates at a higher traffic scale than NYC. For




Figure 3: Forecast quality from the current PostgreSQL-backed result set.

Table 6: Anomaly rates by city

City	Source	Anomaly rate
NYC	mta_station_daily	0.73%
Paris	idfm_station_daily	8.96%

that reason, MAPE is a more meaningful cross-city comparison metric than MAE alone. Even under that normalized view, Paris remains harder to predict. This suggests that future work should either add richer exogenous context or move toward models that capture more complex temporal dependence across the network.

6.3 Anomalies

The anomaly outputs show the strongest contrast between the two systems. `anomaly_rates.csv` reports an anomaly rate of 0.73% for NYC and 8.96% for Paris. This difference is too large to ignore. It suggests that the Paris daily series contains many more observations that look unusual relative to their local and global context under the current anomaly design.

The anomaly table does not by itself explain why those points are unusual, but the distribution is still meaningful. In Paris, major anomalies cluster around strong disruption and recovery periods, especially in the 2020 window. In NYC, the anomaly count is much lower and is concentrated around isolated weak-demand periods such as holidays or atypical weekends. This is precisely the type of screening result the method was designed to produce: it directs attention to time windows that deserve interpretation without claiming causal certainty.

The gap between the two rates is large enough that it should influence how the rest of the report is interpreted. A system with a higher anomaly rate is harder to summarize with one representative weekly pattern and is also more likely to challenge a simple forecasting baseline. In that sense, the anomaly section helps explain the forecasting asymmetry rather than functioning as an isolated appendix.

6.4 Contributors and city structure

The city-level summary confirms that Paris is the larger system in the cleaned reporting scope. From `city_structure_summary.csv`, Paris reaches an average daily demand of approximately 4.13 million, while NYC reaches about 2.63 million. Median daily demand follows the same ordering, and both systems show

Table 7: City structure summary

City	Avg daily	Std. dev.	Wknd./wkday	Peak mo.
NYC	2.63M	1.15M	0.584	2
Paris	4.13M	1.76M	0.512	10

Table 8: Illustrative top contributors

City	Station	Total demand
Paris	SAINT-LAZARE	432.3M
Paris	LA DEFENSE-GRANDE ARCHE	325.3M
Paris	GARE DE LYON	291.6M
NYC	Times Sq-42 St	168.2M
NYC	Grand Central-42 St	115.4M
NYC	34 St-Herald Sq	97.4M

Source: top_contributors.csv.

similar coefficients of variation, around 0.43. The two systems therefore differ more in level and weekly sensitivity than in relative day-to-day volatility.

The weekend-to-weekday ratio makes the weekly difference concrete. NYC keeps a higher weekend share (0.584) than Paris (0.512), which implies that Paris demand is more concentrated in weekday commuting structure. This is consistent with the broader seasonal and temporal narrative from the earlier sections.

The contributor ranking is also plausible after the Stage 2 label cleanup. Paris is dominated by Saint-Lazare, La Défense-Grande Arche, Gare de Lyon, Gare du Nord, and Montparnasse, all of which are obvious interchange or high-throughput nodes. NYC is led by Times Sq-42 St, Grand Central-42 St, 34 St-Herald Sq, and 14 St-Union Sq, which again aligns with domain expectations. The rank order therefore works as a useful validity check on the cleaning pipeline: if the contributor table had been dominated by implausible labels, the rest of the report would have been harder to trust.

The top-contributor view also clarifies why city-level comparisons should not be read as if the two systems were structurally identical. Paris is more top heavy in absolute station totals, while NYC spreads more of its total demand across a larger set of high-volume complexes. That does not invalidate the city comparison; it simply means the comparison is most defensible at the level of cleaned daily totals, weekend sensitivity, and ranked hubs rather than at raw network topology.

6.5 Cross-method synthesis

The four methods are more informative when read together than in isolation. Temporal profiling shows that both cities have strong weekly and seasonal structure, but Paris carries the larger scale and the sharper summer weakening. Forecasting then shows that this larger and less regular Paris signal is harder to predict with the same lag-based baseline. The anomaly layer reinforces that reading by flagging a much higher share of unusual Paris observations. Finally, the contributor ranking shows that Paris demand is more concentrated in a set of very large hubs, while NYC distributes high demand across several major

Table 9: Key quantitative findings

Question	Evidence
Which city has higher average daily demand?	Paris, about 4.13M versus 2.63M for NYC.
Which city is easier to forecast?	NYC, with 5.86% MAPE versus 6.94% for Paris.
Which city shows more unusual days?	Paris, with an 8.96% anomaly rate versus 0.73% for NYC.
What dominates station demand?	Major interchange hubs, led by Saint-Lazare and Times Sq-42 St.

Sources: city_structure_summary.csv, forecast_metrics_by_city.csv, anomaly_rates.csv, and top_contributors.csv.




Figure 4: Anomaly-rate contrast and top contributors in the cleaned result set.

interchange complexes.

This synthesis is one of the main reasons the PostgreSQL-centered contract is useful. Because the daily summaries, anomaly tables, and contributor rankings all come from one cleaned view layer, the project can support a coherent narrative instead of a set of disconnected charts. The final paper therefore depends as much on consistency between methods as on the strength of any single metric.

7 Discussion and Limitations

The final workflow is strong as a reproducible analytical baseline, but it is not a claim that the present method stack is the best available for transport demand modeling. The forecasting component remains a baseline model. It is useful because it is transparent, quick to reproduce, and strong enough to support comparative interpretation, but it does not model richer network or event effects in the way graph-based and more specialized temporal models might [6, 7].

The anomaly layer should also be interpreted carefully. A flagged observation is evidence of unusual behavior, not proof of a particular operational cause. In practice, anomaly detection is most useful here as a triage layer that marks periods for human interpretation or future event matching.

Contributor analysis remains sensitive to source labeling. The project removed mechanical whitespace inconsistencies and restricted Paris ranking to station rows, but semantically related stations can still appear separately when their source identifiers differ. That is acceptable for the current paper because the top-ranked stations remain plausible, but it limits how aggressively the report should interpret small differences among nearby hubs.

Finally, the reproducibility contract depends on gener-




Figure 5: Average daily demand and weekend sensitivity by city.

ated outputs rather than live database queries at paper compile time. This is a reasonable compromise for a project report: the SQL views and workflow define one stable path, and the paper cites the outputs of that path. Still, if the database changes, the workflow must be rerun before the report should be updated.

Despite those limitations, the final project state is substantially stronger than the earlier repository baseline. The notebook no longer carries a separate story from the report, the result tables are regenerated from PostgreSQL rather than assembled manually, and the figures now correspond to current outputs instead of placeholder values. That improvement in analytical discipline is one of the main outcomes of the project.

There is also a broader methodological lesson. For a course project, it is tempting to optimize only the visible final chart or only the most advanced model. The current repository history shows why that is not enough. Once data sources, notebook text, SQL logic, and report claims drift apart, even a good metric becomes hard to defend. The PostgreSQL-first rewrite reduced that risk by forcing all artifacts back onto one result contract.

At the same time, the report should not hide the remaining frontier. Better forecasting likely requires richer external context, more explicit modeling of network dependence, or both. Better anomaly interpretation would require event labeling or an external disruption calendar. More aggressive contributor harmonization would require domain judgment about semantically related stations. Those are meaningful next steps, but they are extensions to a stable baseline, not fixes for a broken workflow.

For planning-oriented interpretation, the current outputs already support a few practical lessons. A city with stronger weekday concentration and higher anomaly rate requires more caution when extrapolating routine patterns into future demand. A system whose dominant stations absorb a large share of total volume may also be more sensitive to local disruptions at those hubs. These are not causal policy recommendations, but they show that the final pipeline can produce operationally meaningful summaries without stepping outside the limits of the available evidence.

8 Conclusion

This project delivers a coherent final analytical baseline for comparing Paris and NYC transport demand. PostgreSQL serves as the single source of truth, reporting-clean views

in the transport schema define the shared contract, and the Python workflow generates the figures and CSV outputs used by the notebook and paper. Within that framework, four methods were enough to produce interpretable findings: temporal profiling established the seasonal and weekly structure, the lag-based forecasting baseline achieved single-digit MAPE in both cities, anomaly detection highlighted much stronger irregularity in Paris, and contributor analysis confirmed plausible dominant stations in both systems.

The most important outcome is not only the set of metrics, but the fact that the project now has a defensible reporting workflow. Each major claim in the paper can be traced back to the cleaned PostgreSQL views and the generated files under `report/results/`. That makes the project suitable for final submission and provides a clear baseline for any future extensions.

References

- [1] X. Ma, Y.-J. Wu, Y. Wang, F. Chen, and J. Liu, “Mining smart card data for transit riders’ travel patterns,” *Transportation Research Part C: Emerging Technologies*, vol. 36, pp. 1–12, 2013.
- [2] Y. Long and J.-C. Thill, “Combining smart card data and household travel survey to analyze jobs–housing relationships in Beijing,” *Computers, Environment and Urban Systems*, vol. 53, pp. 19–35, 2015.
- [3] O. Cats, “Identifying human mobility patterns using smart card data,” arXiv preprint arXiv:2208.05352, 2022.
- [4] N. Aminpour and S. Saidi, “Unveiling mobility patterns beyond home/work activities: A topic modeling approach using transit smart card and land-use data,” *Sustainable Cities and Society*, vol. 114, p. 105745, 2024.
- [5] L. Sun *et al.*, “Optimizing urban mobility through complex network analysis and big data from smart cards,” *IoT*, vol. 6, no. 3, p. 44, 2025.
- [6] P. Xie, M. Ma, T. Li, S. Ji, S. Du, Z. Yu, and J. Zhang, “Spatio-temporal dynamic graph relation learning for urban metro flow prediction,” arXiv preprint arXiv:2204.02650, 2022.
- [7] M. Lablack, H. Qi, Y. Shen, and B. Yin, “ASTIR: Spatio-temporal data mining for crowd flow prediction,” *IEEE Access*, vol. 7, pp. 175159–175165, 2019.
- [8] P. Kolios *et al.*, “Model-adaptive event triggering for monitoring recurrent mobility patterns in public transport,” *IEEE Transactions on Intelligent Transportation Systems*, 2023.
- [9] Y. He, Y. Zhao, and K. L. Tsui, “An analysis of factors influencing metro station ridership: Insights from Taipei Metro,” in *Proc. IEEE International Conference on Intelligent Transportation Systems (ITSC)*, 2018, pp. 1598–1603.

- [10] Y. Chen *et al.*, “A two-stage trip inference model of purposes and socio-economic attributes of regular public transit users,” *IEEE Transactions on Intelligent Transportation Systems*, 2025.
- [11] D. Kundu *et al.*, “Benefits from a new transit line: Exploring the impact of rare users and spatially heterogeneous variations in intensity of use,” *Journal of Public Transportation*, vol. 27, p. 100120, 2025.